

NCAA Championship Statistics



Can we determine March Madness contestants based on historical data?

Overview

The purpose of this analysis was to obtain greater understanding of NCAA Division 1 Statistics using KNN, Random Forest & Logistic Regression

Top Accuracy scores:

- KNN: 66%
- Random Forest: 72%
- Logistic Regression: 75%

Datasets Analyzed

The main dataset used was several data pulls of team stats going back 10 to 20 years, and included such stats as Assist/Turnover Ratio, Three-Point Attempts, and Rebounds, etc. These were all pulled from the stats page of the NCAA website and combined into one main csv file using Spark. The combined data then includes Team Name, Season of the Stats in format of 1011 for example, meaning years 2010 through 2011. Overall, about 3700 records combined to analyze.

Data Cleaning

PROJECT 4

Data

FGP

FGP CSVs

- 01-02 FGP.csv
- 02-03 FGP.csv
- 03-04 FGP.csv
- 04-05 FGP.csv
- 05-06 FGP.csv
- 06-07 FGP.csv
- 07-08 FGP.csv
- 08-09 FGP.csv
- 09-10 FGP.csv
- 10-11 FGP.csv
- 11-12 FGP.csv
- 12-13 FGP.csv
- 13-14 FGP.csv
- 14-15 FGP.csv
- 15-16 FGP.csv
- 16-17 FGP.csv
- 17-18 FGP.csv
- 18-19 FGP.csv
- 19-20 FGP.csv
- 20-21 FGP.csv
- 21-22 FGP.csv
- 22-23 FGP.csv
- 23-24 FGP.csv

FGP XLSX

- FGP 2000-2024.csv

FTP

FTP CSVs

- 01-02 FTP.csv
- 02-03 FTP.csv
- 03-04 FTP.csv
- 04-05 FTP.csv
- 05-06 FTP.csv
- 06-07 FTP.csv
- 07-08 FTP.csv
- 08-09 FTP.csv
- 09-10 FTP.csv
- 10-11 FTP.csv
- 11-12 FTP.csv
- 12-13 FTP.csv
- 13-14 FTP.csv
- 14-15 FTP.csv
- 15-16 FTP.csv
- 16-17 FTP.csv
- 17-18 FTP.csv
- 18-19 FTP.csv
- 19-20 FTP.csv
- 20-21 FTP.csv
- 21-22 FTP.csv
- 22-23 FTP.csv
- 23-24 FTP.csv

	Team	Years	AST	TO	Ratio	TPG	Avg	3FG%	PTS	PPG	ORebs	DRebs	REB	RPG	FGM	FGA	FG%	FT	FTA	FT%
0	Michigan (Big Ten)	12-13	565	368	1.54	9.435897	NaN	38.5	2,932	75.2	0	0	1,366	0.00	1,093	2,260	48.4	450	642	70.1
1	UC Irvine (Big West)	12-13	580	461	1.26	12.459459	NaN	33.4	2,558	69.1	0	0	1,428	0.00	932	2,151	43.3	445	669	66.5
2	Iowa St. (Big 12)	12-13	560	460	1.22	13.142857	NaN	37.4	2,779	79.4	0	0	1,341	0.00	962	2,111	45.6	509	696	73.1
3	Santa Clara (WCC)	12-13	560	468	1.20	12.315789	NaN	36.1	2,827	74.4	0	0	1,315	0.00	970	2,195	44.2	584	803	72.7
4	Wichita St. (MVC)	12-13	530	492	1.08	12.615385	NaN	33.9	2,720	69.7	0	0	1,497	0.00	954	2,166	44.0	553	791	69.9
...
3711	Siena (MAAC)	23-24	342	505	0.68	15.781250	18.8	28.0	1,934	60.4	375	757	1,132	35.38	702	1,758	39.9	361	541	66.7
3712	VMI (SoCon)	23-24	335	507	0.66	15.843750	21.0	31.2	2,204	68.9	299	871	1,170	36.56	818	1,935	42.3	358	517	69.2
3713	Alabama A&M (SWAC)	23-24	369	565	0.65	16.142857	14.7	28.7	2,414	69.0	394	854	1,248	35.66	806	1,918	42.0	654	906	72.2
3714	Mississippi Val. (SWAC)	23-24	271	495	0.55	15.967742	15.8	29.2	1,719	55.5	298	613	911	29.39	613	1,611	38.1	350	532	65.8
3715	Coppin St. (MEAC)	23-24	240	455	0.53	15.689655	18.6	28.3	1,678	57.9	283	590	873	30.10	585	1,571	37.2	355	489	72.6

Problems in training

At first, the goal was to train the model based on individual matchups in an effort to predict a teams performance throughout the NCAA March Madness tournament.

We experienced some issues here, as our data did not include individual matchups so we had to pivot.

We landed on predicting contestants who make it to the tournament, based on season stats.

Contestants who made it into the tournament that year would be given a “Madness” label of 1, and those who did not make it to the tournament would be given a 0.

Resolving Imbalanced Data

Another obstacle we ran into when determining March Madness contestants, is that of 351 Division I NCAA basketball teams, only 68 are considered for the tournament. This presented us with a very imbalanced data set, and in order to fix this imbalance we reduced the number of teams' stats we considered by about half. We kept the top 150 teams by points scored for each season, and then joined the 68 Madness Contestants with the remaining 82 Division I teams.

KNN

- Top accuracy reached: 66%
- Although we tried resampling and scaling our data, KNN was still unable to reach the levels of accuracy we were experiencing with both Random Forest and Logistic Regression
- A couple reasons for this are that KNN is highly sensitive to feature scaling, and imbalanced classes. This causes KNN to struggle in classifying our minority class accurately.
- Finally, when researching optimization, we discovered KNN is sensitive to noisy data.

KNN Classification Report



```
# Print the classification report for the model  
report1 = classification_report(y_test, knn_prediction)  
print(report1)
```

[5]

✓ 0.0s

...

	precision	recall	f1-score	support
0.0	0.79	0.65	0.71	272
1.0	0.50	0.67	0.57	141
accuracy			0.66	413
macro avg	0.64	0.66	0.64	413
weighted avg	0.69	0.66	0.66	413

Random Forest

- Top Accuracy Reached: 72%
- Random Forest was the second classification model we used while trying to optimize Logistic Regression.
- Most optimizations we tried here did not work to improve accuracy and even negatively impacted accuracy in some cases.
- When trying to apply the same optimizations that helped Logistic Regression, we were unsuccessful in reaching the 75% target. This may be because Random Forest is an ensemble of decision trees and could require more tuning and feature engineering for optimal performance.

Random Forest Classification Reports

```
[4] ✓ 0.0s
```

```
# Print classification report 1
print(classification_report(y_test, rf_prediction))
```

...	precision	recall	f1-score	support
0.0	0.79	0.79	0.79	272
1.0	0.59	0.58	0.59	141
accuracy			0.72	413
macro avg	0.69	0.69	0.69	413
weighted avg	0.72	0.72		

```
[13] ✓ 0.0s
```

```
# Print classification report 2
#(After Implementing SelectFromModel Threshold Optimization)
print(classification_report(y_test, rf_prediction))
```

...	precision	recall	f1-score	support
0.0	0.79	0.78	0.78	272
1.0	0.58	0.60	0.59	141
accuracy			0.71	413
macro avg	0.68	0.69	0.68	413
weighted avg	0.72	0.71	0.72	413

```
[16] ✓ 0.0s
```

```
# Print classification report 3
# (After Implementing Class Weight Adjustments)
print(classification_report(y_test, rf_prediction))
```

...	precision	recall	f1-score	support
0.0	0.77	0.79	0.78	272
1.0	0.58	0.55	0.57	141
accuracy			0.71	413
macro avg	0.68	0.67	0.68	413
weighted avg	0.71	0.71	0.71	413

Logistic Regression

- Top Accuracy Reached: 75%
- Logistic Regression was the first model we tried, and originally ran at 74%
- After some “optimizations”, we were not improving the accuracy and this is when we started working on Random Forest and KNN
- Finally, the optimization that got us to our goal, was adjusting the class weights
- We tried adjusting class weights on two different versions, one resampling imbalanced data with SMOTE and one resampling with ADASYN. Both versions came back at 75%. We believe the class weight made the most impact and the resampling had insignificant impact.

Logistic Regression Classification Reports

```
# Classification Report 1
report1 = classification_report(y_test, lr_prediction)
print(report1)
```

[401] ✓ 0.0s

	precision	recall	f1-score	support
0.0	0.75	0.92	0.82	272
1.0	0.72	0.41	0.52	141
accuracy			0.74	413
macro avg	0.73	0.66	0.67	413

```
# Classification Report 3
# SMOTE & Class Weight Adjustments
print(classification_report(y_test, lr_prediction))
```

[408] ✓ 0.0s

	precision	recall	f1-score	support
0.0	0.81	0.81	0.81	272
1.0	0.63	0.64	0.64	141
accuracy			0.75	413
macro avg	0.72	0.72	0.72	413
weighted avg	0.75	0.75	0.75	413

```
# Classification Report 2
# Resampling with SMOTE
print(classification_report(y_test, lr_prediction))
```

[405] ✓ 0.0s

	precision	recall	f1-score	support
0.0	0.82	0.74	0.78	272
1.0	0.58	0.70	0.63	141
accuracy			0.72	413
macro avg	0.70	0.72	0.70	413
weighted avg	0.74	0.72	0.73	413

▶

```
# Classification Report 4
# ADASYN Resampling & Class Weight Adjustments
print(classification_report(y_test, lr_prediction))
```

[417] ✓ 0.0s

	precision	recall	f1-score	support
0.0	0.81	0.80	0.81	272
1.0	0.63	0.65	0.64	141
accuracy			0.75	413
macro avg	0.72	0.72	0.72	413
weighted avg	0.75	0.75	0.75	413

Conclusion

- Had we understood the nature of the problem we wanted to solve a little better at the beginning, we might have spent less time experimenting with incorrect models
- With a more complete data set that included individual season matchups, or more time to put that data together manually, we could have predicted individual matchups
- We are glad our model reached 75% accuracy, however it was clear during optimization that our data is imbalanced. Even at 75% accuracy, our model is predicting teams who will not make the tournament (0s) better than teams that will (1s)